

# Core design patterns for cloud-native architecture

# What is Cloud-Native?

- Cloud-native = apps built to thrive in dynamic, distributed, scalable environments
- Characteristics:
  - Microservices
  - Containers
  - Orchestration
  - CI/CD
  - DevOps and observability

# What are Design Patterns?

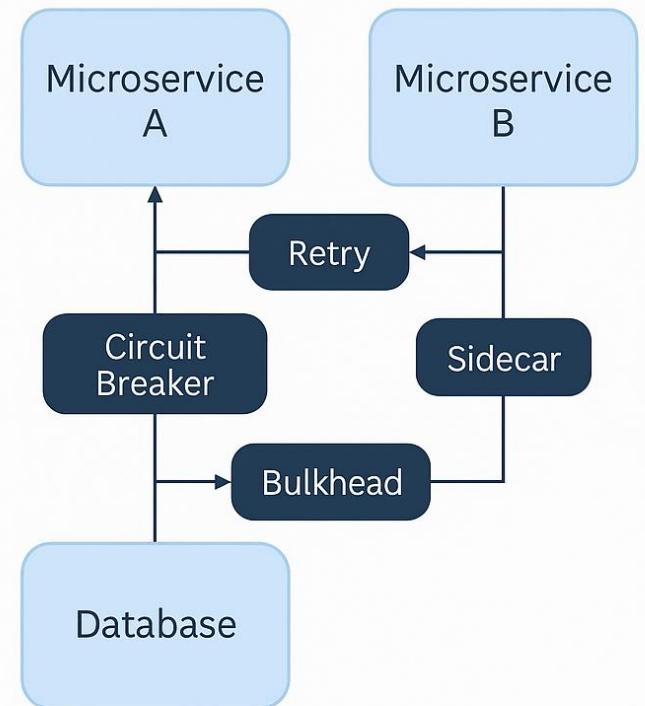
- Reusable templates for solving common software architecture problems.
- Originally popularized in OOP (Gang of Four), but now essential in cloud-native architecture.
- Provide a shared vocabulary for architects and developers.

# Why Design Patterns?

- Patterns = proven solutions to common architectural problems

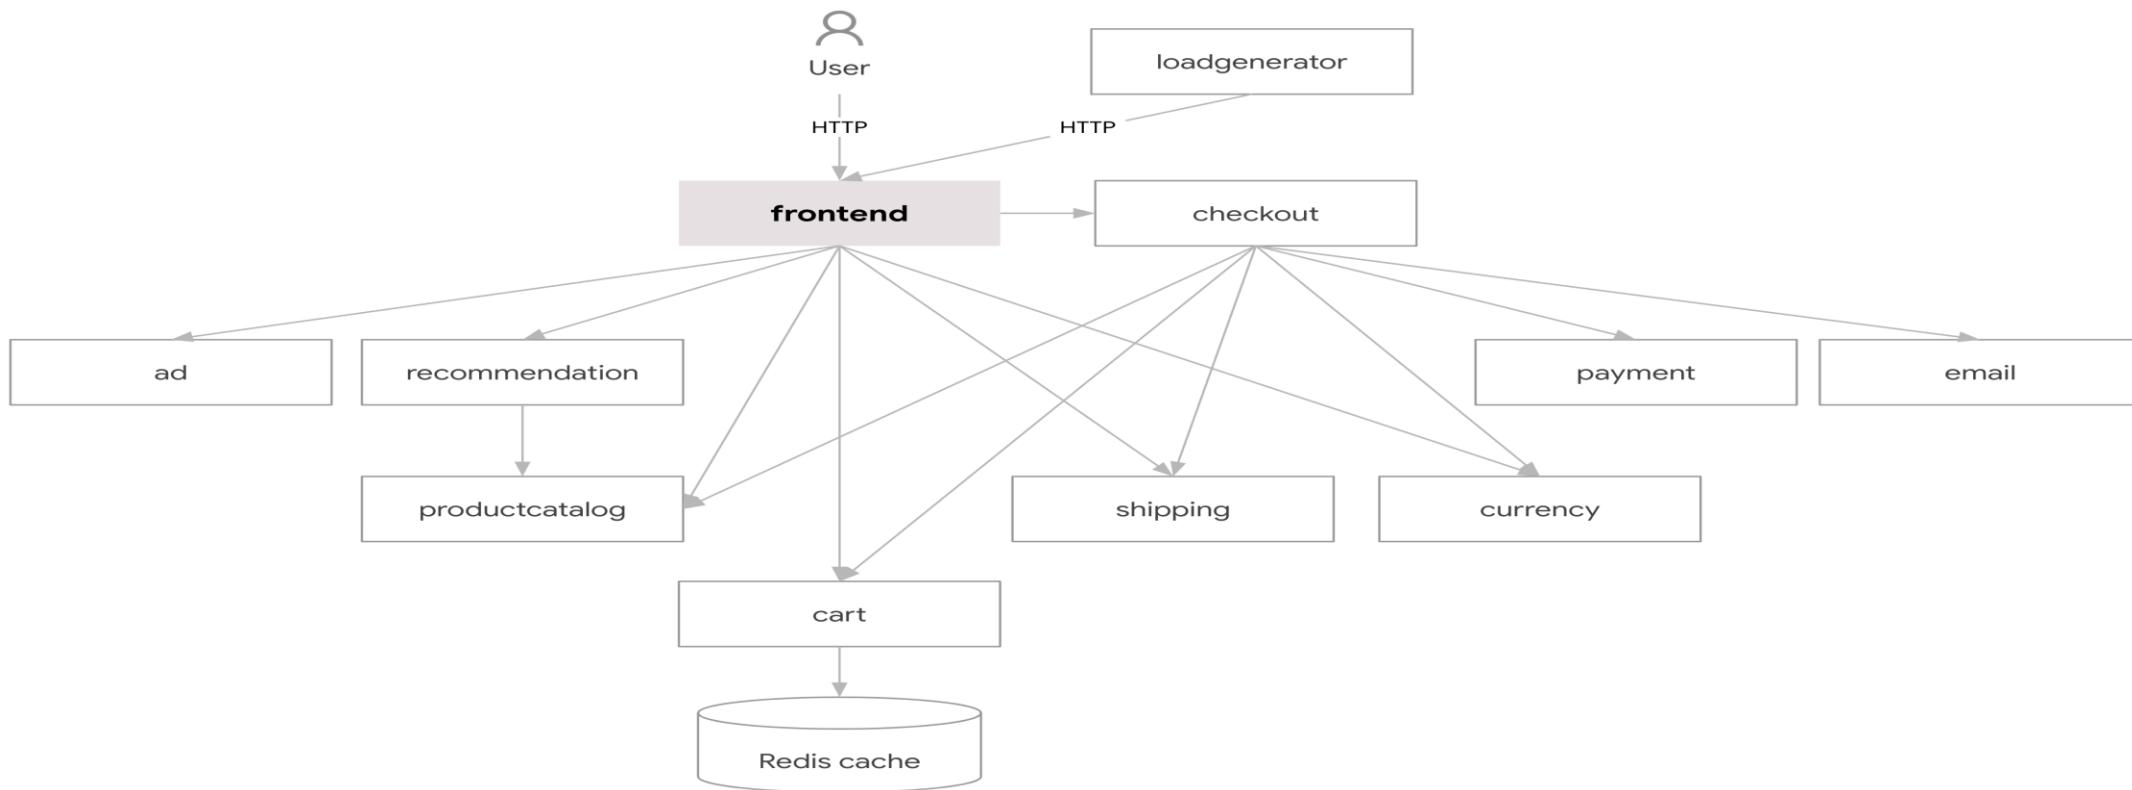
## Why Design Patterns?

- **Consistency:** Patterns ensure uniform solutions across microservices
- **Reusability:** Well-tested patterns can be applied to new problems
- **Scalability:** Patterns help in designing systems that can scale efficiently
- **Resilience:** Make applications fault-tolerant in unreliable environments
- **Observability:** Patterns enable better monitoring and debugging



# Microservice Pattern

- The **Microservice Pattern** breaks down an application into a collection of small, independently deployable services—each responsible for a single business capability.



# Core Principles of Microservices

- Single Responsibility: Each service owns a distinct business function.
- Independent Deployment: Services can be updated without redeploying the entire app.
- Decentralized Data: Each service manages its own database schema.
- Polyglot Technology: Services can use different languages, frameworks, or databases.

# Microservices Architecture

Microservices architecture structures applications as a collection of small, loosely coupled services.

Each service is responsible for a specific business capability.

Benefits:

- Scalability
- Fault Isolation
- Faster Deployment
- Technology Diversity

# Core Patterns in Microservices

- Decomposition Patterns
  - These patterns focus on how to break down an application into services.
- Communication Patterns
  - These define how services interact with each other.
- Data Management Patterns
  - These patterns address how data is handled in a microservices architecture.
- Deployment Patterns
  - These relate to how services are deployed and managed.

# DECOMPOSITION

## Decomposition of Applications According To



# Decomposition Patterns

- By Subdomain
  - This approach follows Domain-Driven Design (DDD) principles.
  - For instance, a logistics company could have subdomains for "Package Tracking" and "Route Optimization."
- By Business Capability
  - Services are structured around business domains.
  - For example, an e-commerce platform might have a "Payment Service," a "User Service," and an "Inventory Service".
- By Strangler Fig Pattern or Strangler Vine Pattern
  - migrating a legacy system to a new system (e.g., microservices architecture) incrementally and safely

# Communication Patterns

- **Synchronous Communication**
  - Synchronous communication is a communication pattern where a service sends a request and waits for a response.
- **Asynchronous Communication**
  - Asynchronous communication allows services to interact without having to communicate sequentially, fostering a loosely coupled, event-driven architecture

# Synchronous Communication

## RESTful APIs

- Example Scenario: A "User Service" fetches user details synchronously using a REST API when a "Notification Service" needs to send alerts.

## gRPC

- Example Scenario: A high-performance "Analytics Service" communicates with a "Data Collection Service" using gRPC.

## CHAINED OR CHAIN OF RESPONSIBILITY

Produces A Single Output Which Is A Combination Of Multiple Chained Outputs.

Use Synchronous HTTP Request Or Response For Messaging



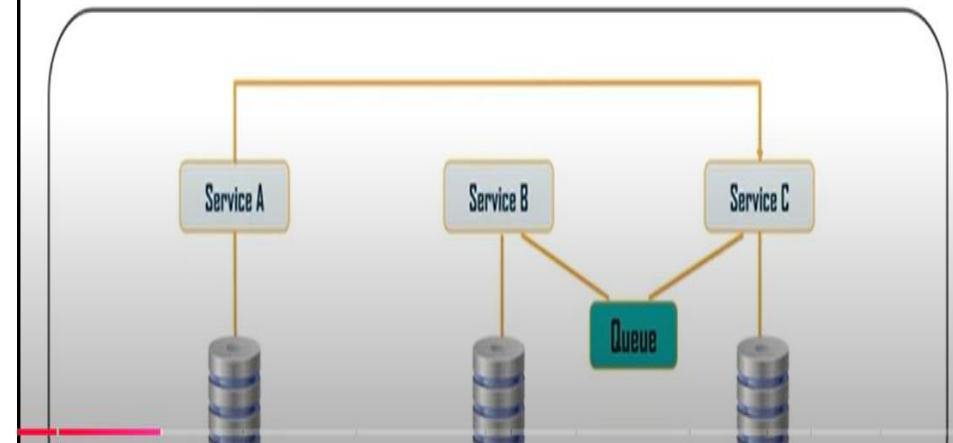
# Asynchronous Communication

## Event Broker

- Event brokers like Kafka or RabbitMQ.
- Example Scenario: An "Order Service" publishes an event when an order is placed, triggering the "Inventory Service" and "Shipping Service" to process the order asynchronously.
- Event-driven architecture for loosely coupled communication

## ASYNCHRONOUS MESSAGING

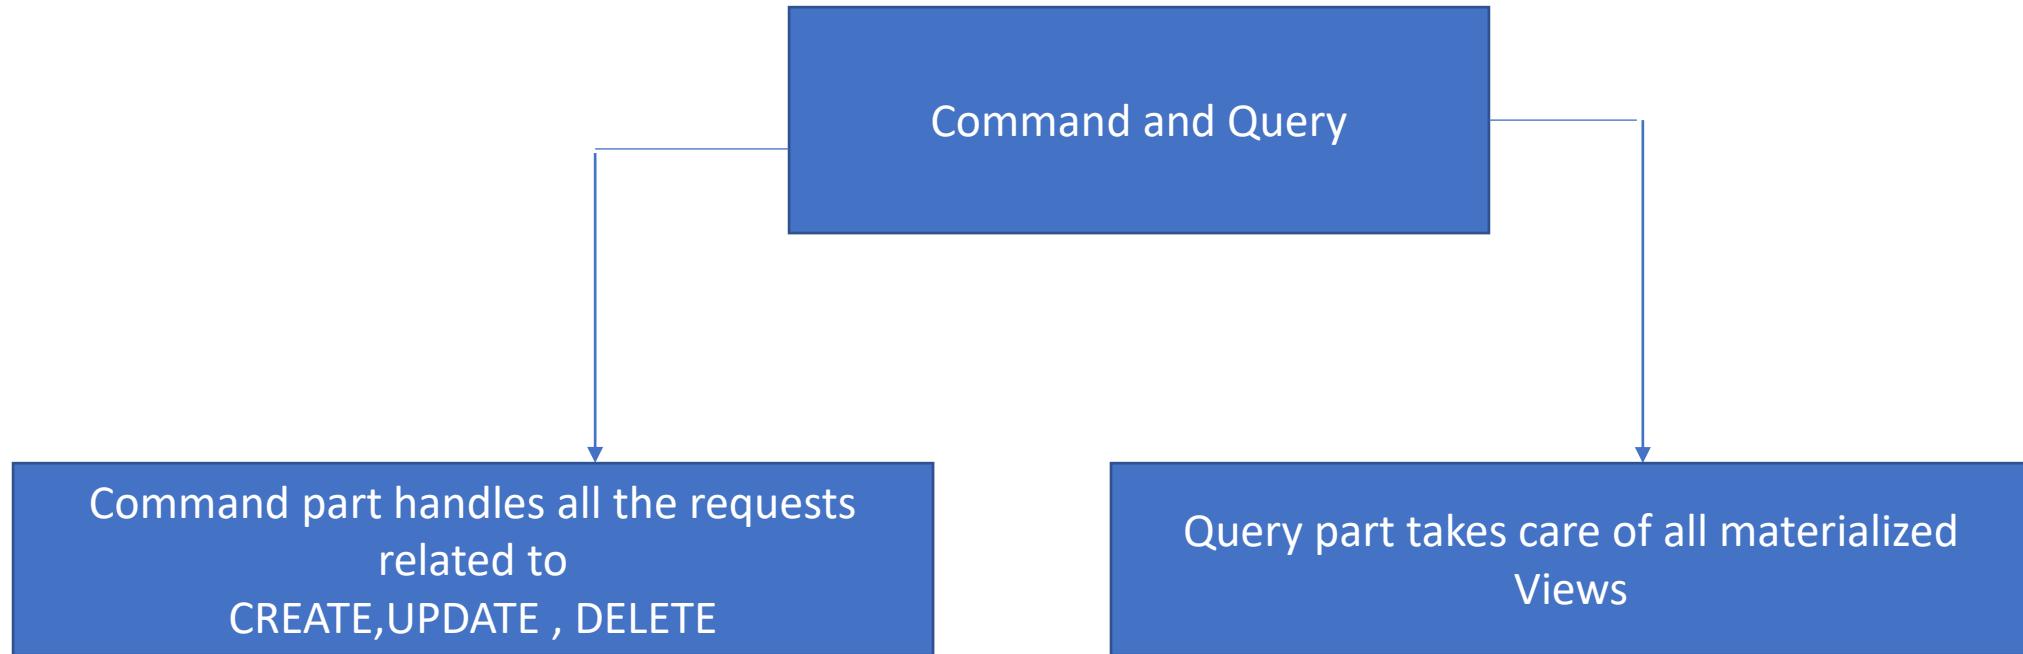
All The Services Can Communicate With Each Other, But They Do Not Have To Communicate With Each Other Sequentially



# Data Management Patterns

- **Database per Service**
  - Ensures independence and avoids tight coupling.
  - Example Scenario: A "Product Service" uses MongoDB for flexibility, while an "Order Service" uses PostgreSQL for transactional consistency.
- **CQRS (Command Query Responsibility Segregation)**
  - Separates read and write operations for better performance.
  - Example Scenario: A "Banking Service" records all account transactions as events to ensure a reliable audit trail.
- **Event Sourcing:**
  - Captures state changes as events.
  - Example Scenario: A "Customer Support Service" uses separate databases for querying customer interactions and updating issue statuses

# CQRS (Command Query Responsibility Segregation)



# CQRS (Command Query Responsibility Segregation)

- **The Scenario: An E-Commerce Product Page**
- Consider a popular product page on an e-commerce website. This page has a very high number of reads compared to writes.
- **Writes (Commands):** Updating the product's price, changing its description, or adding a new unit to inventory. These are infrequent.
- **Reads (Queries):** Thousands of users viewing the product page, its details, stock level, and reviews. These are extremely frequent.

Let's assume a typical read-to-write ratio of **100:1**. For every one time a product's details are updated, it is viewed 10,000 times.

# Traditional Approach (Without CQRS)

In a traditional system, a single database, often normalized for data integrity, handles both reads and writes.

- **Write Operation Cost (C<sub>w</sub>):** A write operation must maintain transactional integrity, potentially updating multiple tables (e.g., Products, Inventory, PriceHistory). This makes it relatively slow.
- Let T<sub>w</sub> (Time for one write) = **80 ms**
- **Read Operation Cost (C<sub>r</sub>):** A read operation on this normalized database must perform complex JOINs to gather all the data for the page (product details, inventory count, review scores, etc.).
- Let T<sub>r</sub> (Time for one read) = **40 ms**

The total computational load (LT) on the single database is the sum of the load from all write and read operations.

$$\begin{aligned} LT &= (1 \times T_w) + (10000 \times T_r) \\ LT &= (1 \times 80 \text{ ms}) + (10000 \times 40 \text{ ms}) \\ LT &= 80 \text{ ms} + 400000 \text{ ms} \\ LT &= 400080 \text{ ms} \end{aligned}$$

In this model, the thousands of expensive read operations create a massive load and contend for resources with the write operations, potentially slowing the entire system down

# CQRS Approach

With CQRS, we have two separate models, each optimized for its task.

**Write Model (Command):** The database is still normalized for data integrity. The cost of a write operation remains the same.

$$Tw-cqrs = \mathbf{80 \text{ ms}}$$

**Read Model (Query):** This is the key difference. The read database is a **denormalized** copy of the data, specifically designed for fast reads. All the data needed for the product page is pre-joined and stored in a single document or a wide table. This makes read operations extremely fast.

$$\text{Let } Tr-cqrs \text{ (Time for one optimized read)} = \mathbf{4 \text{ ms}}$$

Now, the load is split between two independent systems.

**Load on Write System (Lw):**

$$Lw = 1 \times Tw-cqrs = 80 \text{ ms}$$

**Load on Read System (Lr):**

$$Lr = 10000 \times Tr-cqrs = 10000 \times 4 \text{ ms} = 40000 \text{ ms}$$

# Comparison

Metric	Traditional Model	CQRS Model	Performance Gain	
Time per Read	40 ms	4 ms	10x Faster	
Total Read Load	400,000 ms	40,000 ms	10x Less Load	
Scalability	Must scale a single, complex database for both reads and writes.	Can scale the read and write systems independently. Can deploy 10 read-optimized servers and only 1 write server, matching the load perfectly and saving costs.	Highly Optimized	

# Deployment Patterns

- **Single Service per Container:** Simplifies scaling and isolation.
  - Example Scenario: A "Notification Service" runs in its own container, allowing easy scaling during promotional campaigns.
- **Serverless Deployment:** Uses cloud services like AWS Lambda.
  - Example Scenario: An "Image Processing Service" runs as a serverless function to resize and compress images on demand.
- **Orchestrators:** Kubernetes for managing containers.
  - Example Scenario: An "API Gateway" and multiple microservices are managed and scaled using Kubernetes clusters.

# Resilience and Fault-Tolerance Patterns

- **Circuit Breaker**

- Stops repeated failed requests to prevent cascading failures.
- Example Scenario: A "Payment Gateway Service" triggers a circuit breaker if a third-party payment provider becomes unresponsive.

- **Retry**

- Retries failed operations after a delay.
- Example Scenario: A "File Upload Service" retries uploading files to a cloud storage service if there are temporary network issues.

# Resilience and Fault-Tolerance Patterns

- **Bulkhead isolation**

- Allocates resources to prevent overloading critical services.
- Example Scenario: A "Booking Service" isolates resources for "Flight Booking" and "Hotel Booking" to prevent one from impacting the other.

- **Timeout**

- Limits how long a service waits for a response from another service.
- Example Scenario: A "Search Service" sets a timeout for fetching recommendations from an external "Recommendation Service."

# Observability and Monitoring Patterns

- Centralized Logging
  - Aggregates logs using tools like ELK Stack.
  - Example Scenario: A centralized logging system captures logs from "Order Service," "Inventory Service," and "Payment Service" for troubleshooting.
- Distributed Tracing
  - Tools like Jaeger or Zipkin to track requests across services.
  - Example Scenario: Tracing a customer order request across "API Gateway," "Order Service," and "Shipping Service."
- Health Checks
  - Monitors the health of individual services.
  - Example Scenario: An "API Gateway" periodically checks the status of backend services to ensure availability.