

Scaling & Load Balancing in Microservices

Agenda

- Microservices & Scaling Basics
- Vertical vs Horizontal Scaling
- Load Balancing Concepts
- Load Balancing in Microservices
- Tools & Technologies (2025)
- Demo: Node.js with Local Load Balancing
- Summary & Q&A

Introduction to Microservices

- Definition
 - Small, independent services communicating over APIs
- Benefits
 - Scalability, Fault Isolation, Flexibility
- Challenges
 - Distributed nature, network latency, deployment complexity

Why Scaling is Needed

- User base is unpredictable
- Seasonal spikes (e.g., Black Friday, IPL streaming)
- Performance & SLA requirements
- Prevent downtime under heavy load

Amazon Prime Day Traffic Spike

- Prime Day 2023 recorded **375 million items sold globally** in just **48 hours**.
- The event caused a **~4x increase in concurrent traffic** compared to normal days.
- Certain services saw **extreme, uneven load**:
 - Product Search API handled **over 80,000 requests per second** at peak.
 - Checkout & Payment microservices spiked by **6x** during flash sale minutes.
 - Recommendation engine (ML-powered) faced **3x normal load**.

Types of Scaling


- **Vertical Scaling**
 - Add CPU/RAM to same machine
 - Pros: Simple, no code changes
 - Cons: Hardware limits, single point of failure
- **Horizontal Scaling**
 - Add more instances/machines
 - Pros: Fault-tolerant, cheaper long term
 - Cons: Needs load balancing, distributed stat

Auto Scaling in Microservices

- Cloud-native auto scaling services
 - AWS Auto Scaling Groups
 - Azure Virtual Machine Scale Sets
 - Kubernetes Horizontal Pod Autoscaler (HPA)
- Triggered by metrics
 - CPU usage
 - Request count
 - Custom business metrics

Load Balancing Basics

- Definition: Distributing traffic across multiple servers
- Goals: Improve responsiveness, availability, resource utilization
- Strategies: Round Robin, Least Connections, IP Hash



Strategies

- Round Robin – Equal rotation
- Least Connections – Send to least busy server
- IP Hash – Consistent routing based on client IP
- Weighted Round Robin – More traffic to powerful servers

Load Balancing in Microservices

- Service discovery integration
- Dynamic instance registration/deregistration
- Examples:
 - Nginx,
 - HAProxy,
 - Envoy,
 - Kubernetes Ingress

Service Discovery Tools

- Consul: Key-value store, health checks, DNS interface
- Eureka: Netflix OSS service registry
- Kubernetes DNS: Internal service discovery

Scaling in Kubernetes

- Horizontal Pod Auto-scaler (HPA)
- Cluster Auto-scaler
- Resource Requests & Limits

Cloud-native Scaling Tools

- AWS Auto Scaling Groups
- GCP Instance Groups
- Azure VM Scale Sets


Load Balancing in Cloud Platforms

- AWS Elastic Load Balancer (ALB/NLB)
- GCP Load Balancer
- Azure Load Balancer

Challenges in Scaling

- Data consistency
- Stateful services
- Cost optimization

Challenges of Scaling Microservices



Cost Optimization



Best Practices

- Stateless service design
- Monitoring & alerting
- Graceful degradation

Demo Plan

- Node.js microservice with frontend
- Run multiple instances
- Configure Nginx load balancer
- Test with load testing tool (e.g., Apache JMeter)

Demo Architecture

- Frontend: React
- Backend: Node.js + Express
- Load Balancer: Nginx
- Service Discovery: Consul

Demo Step-by-Step

- Build Node.js service
- Create Dockerfile
- Run multiple containers
- Configure Nginx for round-robin
- Verify scaling

Latest Tech Stack

- Node.js 20+
- React 18+
- Docker 25+
- Consul 1.16+
- Nginx 1.25+

Monitoring Tools

- Prometheus + Grafana
- ELK Stack
- Datadog

Feature	Scaling	Load Balancing
Definition	Increasing or decreasing the capacity of your system to handle load.	Distributing incoming requests evenly across multiple instances or servers.
Goal	To handle more requests by adding computing power (scale up) or more instances (scale out).	To ensure no single instance is overwhelmed and traffic is handled efficiently.
Types	<ul style="list-style-type: none"> - Vertical Scaling (Scale-Up): Adding more CPU/RAM to an instance. - Horizontal Scaling (Scale-Out): Adding more instances of a service. 	<ul style="list-style-type: none"> - DNS Load Balancing (Round-robin, weighted). - Application Load Balancer (Layer 7). - Network Load Balancer (Layer 4).
When Needed	When your current system cannot handle peak traffic or performance degrades.	When you already have multiple instances running and need to split traffic.

Feature	Scaling	Load Balancing
Example in Microservices	Adding 5 more pods for the payment-service in Kubernetes when CPU > 80%.	Routing checkout requests evenly to all payment-service pods so no single pod is overloaded.
Cloud/DevOps Tools	- Kubernetes Horizontal Pod Autoscaler (HPA) - AWS Auto Scaling Groups - Azure VM Scale Sets - Google Cloud Instance Groups	- AWS Elastic Load Balancing (ELB) - Nginx / HAProxy - Kubernetes Ingress Controller
Key Metrics	CPU/Memory utilization, request rate (RPS), queue length, response time.	Response time per instance, request distribution, health checks.
Relation	Scaling adds capacity ; load balancing manages traffic across capacity .	Load balancing needs capacity ; scaling provides that capacity .

Conclusion

- Scaling & load balancing are key for performance
- Choose right tools for your architecture
- Automate wherever possible

Q&A

- Open discussion